

# Machine learning foundations I

Mitko Veta

Eindhoven University of Technology  
Department of Biomedical Engineering

2020

# Learning goals

At the end of this lecture you will:

- ▶ Have an understanding of the goal of machine learning (ML) models.
- ▶ Have a good understanding of basic mathematical concepts used in ML and be able to apply them in the design and implementation of ML methods.

# Overview

Topics covered in this lecture:

1. Linear algebra
2. Gradient-based optimization
3. Two simple machine learning models
  - Linear model
  - Nearest-neighbours model
4. Probability theory

# Note on the slides

This set of slides is larger than the one used during the lectures. It includes some additional material that you can use as a guide when studying.

# Linear algebra

Materials:

- ▶ Chapter 1.2 from **deeplearning**
- ▶ **linearalgebra**

# Scalars

- ▶ A scalar is a single number (integer, real, rational, ...).
- ▶ Denoted by italics  $a, n, x$

# Vectors

- ▶ A vector is a 1-D array of numbers (integer, real, rational, ...)

$$\mathbf{x} = \begin{bmatrix} x_1 \\ x_2 \\ \dots \\ x_n \end{bmatrix}$$

- ▶ Example notation for type and size  
 $\mathbf{x} \in \mathbb{R}^n$

# Matrices

- ▶ A matrix is a 2-D array of numbers

$$\begin{bmatrix} a_{1,1} & a_{1,2} \\ a_{2,1} & a_{2,2} \end{bmatrix}$$

- ▶ Example notation for type and shape

$$\mathbf{A} \in \mathbb{R}^{m \times n}$$



# Tensors

- ▶ A tensor is an array of numbers that may have
  - ▶ a zero dimensions and be a scalar,
  - ▶ one dimension and be a vector,
  - ▶ two dimensions and be a matrix,
  - ▶ more dimensions ...

**Side note:** One of the most popular frameworks for implementing deep machine learning models is called TensorFlow (<https://www.tensorflow.org/>).

# Transpose matrix

$$(\mathbf{A}^T)_{i,j} = \mathbf{A}_{j,i}$$

$$\mathbf{A} = \begin{bmatrix} A_{1,1} & A_{1,2} \\ A_{2,1} & A_{2,2} \\ A_{3,1} & A_{3,2} \end{bmatrix} \Rightarrow \begin{bmatrix} A_{1,1} & A_{2,1} & A_{3,1} \\ A_{1,2} & A_{2,2} & A_{3,2} \end{bmatrix}$$

The transpose matrix is a mirror image with regard to the main diagonal

$$(\mathbf{AB})^T = \mathbf{B}^T \mathbf{A}^T$$

# Identity matrix

- ▶ Identity matrix  $I_3$

$$\begin{bmatrix} 1 & 0 & 0 \\ 0 & 1 & 0 \\ 0 & 0 & 1 \end{bmatrix}$$

- ▶ The identity matrices are neutral elements in matrix-matrix and matrix-vector multiplication, e.g.

$$\forall \mathbf{x} \in \mathbb{R}^n : I_n \mathbf{x} = \mathbf{x} I_n = \mathbf{x}$$

## Matrix (dot) product

$$\mathbf{C} = \mathbf{AB}$$

The matrices must be compatible: an  $m \times n$  matrix is multiplied with an  $n \times r$  matrix and as a result an  $m \times r$  matrix is obtained

$$C_{i,j} = \sum_k A_{i,k} B_{k,j}$$

$$\mathbf{A} \begin{bmatrix} 1 & 2 & 3 \\ 4 & 5 & 6 \\ 7 & 8 & 9 \\ 10 & 11 & 12 \end{bmatrix} \times \mathbf{B} \begin{bmatrix} 1 & 2 & 3 & 4 & 5 \\ 6 & 7 & 8 & 9 & 10 \\ 11 & 12 & 13 & 14 & 15 \end{bmatrix} = \mathbf{C} [4 \times 5]$$

$$C_{2,5} = A_{2,1}B_{1,5} + A_{2,2}B_{2,5} + A_{2,3}B_{3,5} = 4 \cdot 5 + 5 \cdot 10 + 6 \cdot 15 = 160$$

# Matrix (dot) product

- ▶ In general matrix multiplication is not commutative, i.e., most of the time  $\mathbf{AB} \neq \mathbf{BA}$ .
- ▶ Depending on the dimensions sometimes  $\mathbf{AB}$  or  $\mathbf{BA}$  are not possible.
- ▶ As a special case the matrix can be a (column or row) vector; an  $m \times n$  matrix is multiplied with a  $n \times 1$  vector to obtain a  $m \times 1$  vector.

# Systems of linear equations

$$A_{1,1}x_1 + A_{1,2}x_2 + \dots + A_{1,n}x_n = b_1$$

$$A_{2,1}x_1 + A_{2,2}x_2 + \dots + A_{2,n}x_n = b_2$$

...

$$A_{m,1}x_1 + A_{m,2}x_2 + \dots + A_{m,n}x_n = b_m$$

- ▶  $A_{*,*}$  and  $b_*$  are the knowns,  $x_*$  are the unknowns.
- ▶ In matrix form:  **$A\mathbf{x} = \mathbf{b}$**

# Systems of linear equations

►  $\mathbf{Ax} = \mathbf{b}$  expands to

$$\mathbf{A}_{1,:}\mathbf{x}_1 = \mathbf{b}_1$$

$$\mathbf{A}_{2,:}\mathbf{x}_2 = \mathbf{b}_2$$

...

$$\mathbf{A}_{m,:}\mathbf{x}_m = \mathbf{b}_m$$

# Solving systems of linear equations

- ▶ A linear system of equations can have
  - ▶ no solutions,
  - ▶ many solutions,
  - ▶ exactly one solution.
- ▶ Only one solution implies that multiplication by a matrix is an invertible operation.



# Matrix inversion

- ▶ Matrix inverse is defined with

$$\mathbf{A}^{-1}\mathbf{A} = \mathbf{I}_n$$

- ▶ A system of linear equations can be solved using inverse matrix

$$\mathbf{Ax} = \mathbf{b}$$

$$\mathbf{A}^{-1}\mathbf{Ax} = \mathbf{A}^{-1}\mathbf{b}$$

$$\mathbf{I}_n\mathbf{x} = \mathbf{A}^{-1}\mathbf{b}$$

$$\mathbf{x} = \mathbf{A}^{-1}\mathbf{b}$$

- ▶ This is useful mostly for abstract analysis.
- ▶ From a numerical point of view there are much more efficient methods.

# Invertibility

A matrix cannot be inverted if

- ▶ the number of rows and columns is not the same, or
- ▶ some rows and columns are "redundant" ("linearly dependent", "low rank").

# Moore-Penrose pseudoinverse

- ▶ Matrix inversion is not defined on matrices that are not square.
- ▶ The **Moore-Penrose pseudoinverse** is defined as

$$\mathbf{A}^+ = \lim_{\alpha \searrow 0} (\mathbf{A}^T \mathbf{A} + \alpha \mathbf{I})^{-1} \mathbf{A}^T$$

# Moore-Penrose pseudoinverse

Now we can consider

$$\mathbf{x} = \mathbf{A}^+ \mathbf{y}$$

- ▶ If the equation has
  - ▶ exactly one solution: this is the same as inverse,
  - ▶ no solution: gives the solution with the smallest error,  $\|\mathbf{Ax} - \mathbf{y}\|_2$
  - ▶ many solutions: gives the solution with the smallest norm of  $\mathbf{x}$ .

# Singular value decomposition

- ▶ Similar to eigenvalue decomposition
- ▶ More general: matrix need not be square

$$\mathbf{A} = \mathbf{U}\mathbf{D}\mathbf{V}^T$$

- ▶  $\mathbf{U}$  and  $\mathbf{V}$  are square matrices and are both orthogonal,  $\mathbf{D}$  is diagonal.
- ▶ The diagonal elements of  $\mathbf{D}$  are called **singular values** of matrix  $\mathbf{A}$ ; the columns of  $\mathbf{U}$  and  $\mathbf{V}$  are **left-singular** and **right-singular vectors** of  $\mathbf{A}$ , respectively.

# Computing the pseudoinverse

- ▶ Efficient implementations are based on the formula allowed by the singular decomposition

$$\mathbf{A}^+ = \mathbf{V}\mathbf{D}^+ \mathbf{U}^T$$

- ▶  $\mathbf{U}$ ,  $\mathbf{D}$ ,  $\mathbf{V}$  are from the singular value decomposition of  $\mathbf{A}$ .
- ▶ The pseudoinverse  $\mathbf{D}^+$  of  $\mathbf{D}$  is obtained by taking the reciprocal non-zero elements and after that taking the transpose of the resulting matrix.

# Norms

- ▶ Norms are functions that measure how "large" a vector is.
- ▶ Similar to a distance between zero and the point represented by the vector
  - ▶  $f(\mathbf{x}) = 0 \Rightarrow \mathbf{x} = 0$
  - ▶  $f(\mathbf{x} + \mathbf{y}) \leq f(\mathbf{x}) + f(\mathbf{y})$  (**the triangle inequality**)
  - ▶  $\forall \alpha \in \mathbb{R} : f(\alpha \mathbf{x}) = \alpha f(\mathbf{x})$

# Norms

- ▶  $L^p$ - norm

$$\|\mathbf{x}\|_p = \left( \sum_i |x_i|^p \right)^{\frac{1}{p}}$$

- ▶ Most popular  $L^2$ -norm (for  $p = 2$ )
- ▶  $L_1$ -norm (for  $p = 1$ ):  $\|\mathbf{x}\|_1 = \sum_i |x_i|$
- ▶ Max norm (for infinite  $p$ ):  $\|\mathbf{x}\|_\infty = \max_i |x_i|$



# Special vectors and matrices

- ▶ Unit vector  $\|\mathbf{x}\|_n = 1$
- ▶ Symmetric matrix  $\mathbf{A} = \mathbf{A}^T$
- ▶ Orthogonal matrix

$$\mathbf{A}\mathbf{A}^T = \mathbf{I} = \mathbf{A}^T\mathbf{A}$$

- ▶ It follows that for orthogonal matrices  $\mathbf{A}^T = \mathbf{A}^{-1}$

# Eigendecomposition

- ▶ Eigenvector and eigenvalue

$$\mathbf{A}\mathbf{x} = \lambda\mathbf{x}$$

- ▶ Eigendecomposition of a matrix

$$\mathbf{A} = \mathbf{V}\text{diag}(\lambda)\mathbf{A}^{-1}$$

where  $\text{diag}(\lambda)$  is a diagonal matrix having the (scalar) eigenvalues  $\lambda$  as diagonal elements.

# Eigendecomposition

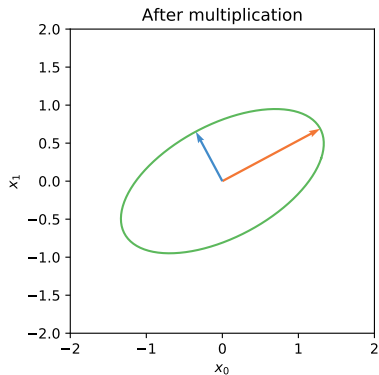
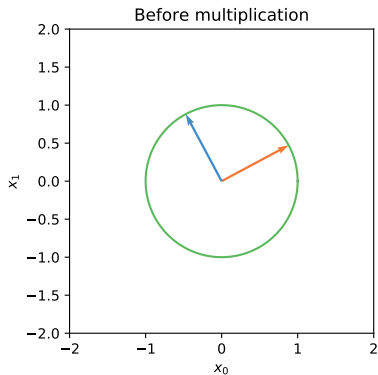
- ▶ Every real symmetric matrix has a real orthogonal eigendecomposition

$$\mathbf{A} = \mathbf{Q}\mathbf{\Lambda}\mathbf{Q}^T$$

where  $\mathbf{Q}$  is an orthogonal matrix composed of eigenvectors of  $\mathbf{A}$  and  $\mathbf{\Lambda}$  is a diagonal matrix.

- ▶ The eigenvalue  $\Lambda_{ii}$  is associated with the eigenvector in column  $i$  of  $\mathbf{Q}$ , denoted as  $\mathbf{Q}_{:,i}$ .
- ▶ We can think of  $\mathbf{A}$  as scaling space by factor  $\lambda_i$  in the direction of its corresponding eigenvector  $\mathbf{v}^{(i)}$  (represented by  $\mathbf{Q}_{:,i}$ ).

# Effect of eigenvalues



# Eigendecomposition

- ▶ From the eigendecomposition we learn useful properties of the matrix.
- ▶ The eigendecomposition of a real symmetric matrix is used in optimization of quadratic expressions of the form  $f(\mathbf{x}) = \mathbf{x}^T \mathbf{A} \mathbf{x}$  under the constraint  $\|\mathbf{x}\|_2 = 1$ .
- ▶ For instance, if  $\mathbf{x} = \mathbf{v}^{(i)}$ , then  $f(\mathbf{x}) = \lambda_i$ , when  $\mathbf{v}^{(i)}$  is an eigenvector of  $A$  and  $\lambda_i$  is its corresponding eigenvalue.
- ▶ The maximal (minimal) value of  $f$  within the constraint region is equal to the maximal (minimal) eigenvalue.

- ▶ A **trace** of a matrix is defined as

$$Tr(\mathbf{A}) = \sum_i \mathbf{A}_{i,i}$$

- ▶ Expressions in terms of the trace operators allow to exploit many useful identities, e.g.

$$Tr(\mathbf{ABC}) = Tr(\mathbf{BCA}) = Tr(\mathbf{CAB})$$

# Gradient-based optimization

Materials:

- ▶ Chapters 1.4 and 1.5 from **deeplearning**
- ▶ **linearalgebra**

# Gradient

- ▶ Let  $f : \mathbb{R}^{m \times n} \mapsto \mathbb{R}$  be a function that takes  $m \times n$  matrix  $\mathbf{A}$  as input and returns a real number (scalar).
- ▶ A **gradient** of  $f$  with respect to  $A$  is the matrix

$$\nabla_{\mathbf{A}} f(\mathbf{A}) = \begin{bmatrix} \frac{\partial f}{\partial A_{11}} & \frac{\partial f}{\partial A_{12}} & \cdots & \frac{\partial f}{\partial A_{1n}} \\ \frac{\partial f}{\partial A_{21}} & \frac{\partial f}{\partial A_{22}} & \cdots & \frac{\partial f}{\partial A_{2n}} \\ \vdots & \vdots & \ddots & \vdots \\ \frac{\partial f}{\partial A_{m1}} & \frac{\partial f}{\partial A_{m2}} & \cdots & \frac{\partial f}{\partial A_{mn}} \end{bmatrix}$$

- ▶ i.e. an  $m \times n$  matrix with

$$(\nabla_{\mathbf{A}} f(\mathbf{A}))_{ij} = \frac{\partial f}{\partial A_{ij}}$$

- ▶ The size of the gradient of  $\mathbf{A}$  is the same as the size of  $A$ .



# Gradient

- ▶ In the special case when  $A$  is a vector we obtain the (possibly more familiar) gradient

$$\nabla_{\mathbf{x}} f(\mathbf{x}) = \begin{bmatrix} \frac{\partial f}{\partial x_1} \\ \frac{\partial f}{\partial x_2} \\ \vdots \\ \frac{\partial f}{\partial x_m} \end{bmatrix}$$

- ▶ In general to define a gradient we require that the function returns a **real** value.

# Jacobian

- ▶ The Jacobian  $\mathbf{J}_f$  is a generalization of the gradient for vector valued functions.
- ▶ Let  $\mathbf{f} : \mathbb{R}^n \mapsto \mathbb{R}^m$  be a function that takes  $n$ -dimensional vector  $\mathbf{x}$  as input and returns a  $m$ -dimensional vector as an output.
- ▶ The Jacobian  $\mathbf{J}_f$  is defined as

$$\mathbf{J}_f = \begin{bmatrix} \frac{\partial f_1}{\partial x_1} & \frac{\partial f_1}{\partial x_2} & \cdots & \frac{\partial f_1}{\partial x_n} \\ \frac{\partial f_2}{\partial x_1} & \frac{\partial f_2}{\partial x_2} & \cdots & \frac{\partial f_2}{\partial x_n} \\ \vdots & \vdots & \ddots & \vdots \\ \frac{\partial f_m}{\partial x_1} & \frac{\partial f_m}{\partial x_2} & \cdots & \frac{\partial f_m}{\partial x_n} \end{bmatrix}$$

- ▶ Note that for the special case of a scalar-valued function, the Jacobian is the transpose of the gradient.

# Optimization

- ▶ Most machine learning methods involve some kind of optimization.
  - ▶ One exception is the  $k$ -Nearest neighbour classifier introduced later.
- ▶ Optimization means minimizing or maximizing some function  $f(\mathbf{x})$ , i.e. finding the values of  $\mathbf{x}$  for which  $f(\mathbf{x})$  has a minimum or a maximum.
- ▶ Notation:  $\mathbf{x}^* = \operatorname{argmin} f(\mathbf{x})$

# Gradient-based optimization

- ▶ The derivative tells us how to change  $x$  in order to make a small improvement of  $f(x)$ .
- ▶ Therefore, derivatives can be useful in optimization.

# Two simple machine learning models

Materials:

- ▶ Chapter 2.3 from **elements**

## Some notations

- ▶ We denote an input variable with the symbol  $x$  (scalar) or  $\mathbf{x}$  (vector).
- ▶ The  $i$ -th component of a vector input  $\mathbf{x}$  is denoted as  $x_i$ .
- ▶ Quantitative (numerical) outputs are denoted with  $y$ .
- ▶ Qualitative outputs are denoted with  $g$  (from group) and take values from a set  $\mathcal{G}$ .
- ▶ Matrices are denoted with bold and uppercase letters  $\mathbf{X}$  for instance, a set of  $N$  input  $p$ -vectors  $\mathbf{x}_i$  ( $1 \leq i \leq N$ ) is "packed" in a  $N \times p$  input matrix  $\mathbf{X}$ .
- ▶ Since by default vectors are assumed to be column vectors, the rows of  $\mathbf{X}$  are the transposes  $\mathbf{x}_i^T$ .

# The learning task

- ▶ Given a value of the input vector  $\mathbf{x}$  make a good prediction of the output  $y$ , denoted as  $\hat{y}$ .
- ▶ Both  $y$  and  $\hat{y}$  should take values from the same numerical set.
- ▶ Similarly,  $g$  and  $\hat{g}$  should both take values from the same set  $\mathcal{G}$ .
- ▶ We suppose that we have available a set of measurements  $(\mathbf{x}_i, y_i)$  or  $(\mathbf{x}_i, g_i)$  ( $1 \leq i \leq N$ ) called **training data** (in matrix form:  $(\mathbf{X}, \mathbf{y})$  and/or  $(\mathbf{X}, \mathbf{g})$ ).
- ▶ Our task is to construct a prediction rule based on the training data.

# The learning task

Example:

- ▶ **Variable values:** Let  $g$  (and therefore also  $\hat{g}$ ) be two valued (categorical), e.g.  $\mathcal{G} = \{\text{BLUE}, \text{ORANGE}\}$ .
- ▶ **Encoding of  $g$ s with  $y$ s:** Then each class can be encoded binary, i.e., with  $y \in \{0, 1\}$ , e.g., **BLUE** and **ORANGE**, would correspond to 0 and 1, respectively.
- ▶ **Predicted output values:**  $\hat{y}$  ranges over the interval  $[-\infty, +\infty]$  (of which  $\{0, 1\}$  is a subset).
- ▶ **Prediction rule:**  $\hat{g}$  is assigned a (class label) **BLUE** if  $\hat{y} < 0.5$  and **ORANGE**, otherwise.



# Two simple approaches to prediction

- ▶ Linear model fit
  - ▶ strong assumptions about the structure of the decision boundary
- ▶  $k$ -nearest neighbours
  - ▶ weak assumptions about the structure of the decision boundary

# Linear model fit by least squares

- ▶ Despite relative simplicity one of the most important statistical tools
- ▶ Input vector  $\mathbf{x}^T = (x_1, x_2, \dots, x_p)$
- ▶ Output  $y$  predicted using the model

$$\hat{y} = \hat{w}_0 + \sum_{j=1}^p x_j \hat{w}_j$$

- ▶  $\hat{w}_i$  ( $0 \leq i \leq p$ ) are the parameters of the linear model
- ▶ In vector form

$$\hat{y} = \hat{\mathbf{w}}^T \mathbf{x} = \mathbf{x}^T \hat{\mathbf{w}}$$

using the fact that the scalar (inner) product of two vectors is a commutative operation.

# Linear model fit by least squares

- ▶ We assume that  $w_0$  is in  $\mathbf{w}$  and 1 is included in  $\mathbf{x}$ .
- ▶  $\hat{y}$  is a scalar, but in general can be a  $k$ -vector  $\hat{\mathbf{y}}$ , in which case  $\mathbf{w}$  becomes a  $p \times k$  matrix of coefficients.

# Linear model fit by least squares

Some hyper(space) terminology:

- ▶ Points  $\mathbf{x}, \hat{y}$  form a **hyperplane** in the  $(p + 1)$ -dimensional input-output hyperspace.
- ▶ If  $\mathbf{x}$  is extended with constant 1 then the hyperplane includes the origin and it forms a **subspace**.
- ▶ If 1 is not included then the hyperplane is an **affine** set and it cuts the  $y$ -axis at the point  $(\mathbf{0}, \hat{w}_0)$ , where the vector  $\mathbf{0}$  has all  $x_i$  coordinates equal to 0.
- ▶ Reminder: from now on we assume that 1 is included in  $\mathbf{x}$  and  $\hat{w}_0$  in  $\hat{\mathbf{w}}$ .
- ▶ The function  $f(\mathbf{x}) = \mathbf{w}^T \mathbf{x}$  defined on the  $p$ -dimensional (input) space is a **linear** function (we omit the hats over the  $\mathbf{w}$ s since now we consider them as free variables).
- ▶ The gradient  $\nabla f(\mathbf{x})$  is a vector pointing along the direction of maximal change.

# Linear model fit by least squares

- ▶ There are many ways to fit a linear model to a training dataset.
- ▶ **Least squares** method
  - ▶ We need to find coefficients  $\hat{w}_i$  which minimize the error estimated with the **residual sum of squares**

$$\text{RSS}(\mathbf{w}) = \sum_{i=1}^N (y_i - \mathbf{x}_i^T \mathbf{w})^2$$

assuming  $N$  input-output pairs.

- ▶  $\text{RSS}(\mathbf{w})$  is a quadratic function.
- ▶ A minimum always exists though not necessarily a unique one.

# Linear model fit by least squares

- ▶ We look for the solution  $\hat{\mathbf{w}}$  using the matrix notation:
- ▶  $\mathbf{y} = [y_1, y_2, \dots, y_N]^T$  is the vector formed from the  $N$  output vectors and  $\mathbf{X}$  is an  $N \times p$  matrix

$$\text{RSS}(\mathbf{w}) = (\mathbf{y} - \mathbf{X}\mathbf{w})^T(\mathbf{y} - \mathbf{X}\mathbf{w})$$

- ▶ To find the minimum we differentiate with respect to  $\mathbf{w}$  which gives

$$-2\mathbf{X}^T(\mathbf{y} - \mathbf{X}\mathbf{w})$$

For details about the derivation check equations 4 and 5 in [this document](#).

# Linear model fit by least squares

- ▶ To find the minimum our derivative must be **0**, hence:

$$\mathbf{X}^T(\mathbf{y} - \mathbf{X}\mathbf{w}) = \mathbf{0}$$

$$\mathbf{X}^T\mathbf{y} - \mathbf{X}^T\mathbf{X}\mathbf{w} = \mathbf{0}$$

$$\mathbf{X}^T\mathbf{y} = \mathbf{X}^T\mathbf{X}\mathbf{w}$$

- ▶ If  $\mathbf{X}^T\mathbf{X}$  is non-singular there exists a unique solution given by

$$\hat{\mathbf{w}} = (\mathbf{X}^T\mathbf{X})^{-1}\mathbf{X}^T\mathbf{y}$$

## Discussion point

Why we cannot simply solve for  $\hat{\mathbf{w}}$  in the following way?

$$\mathbf{y} - \mathbf{X}\mathbf{w} = \mathbf{0}$$

$$\mathbf{y} = \mathbf{X}\mathbf{w}$$

$$\hat{\mathbf{w}} = \mathbf{X}^{-1}\mathbf{y}$$



# Linear model fit by least squares

- ▶ For each input  $\mathbf{x}_i$  there corresponds the fitted output

$$\hat{y}_i = \hat{y}_i(\mathbf{x}_i) = \hat{\mathbf{w}}^T \mathbf{x}_i$$

.

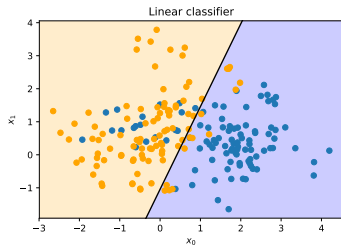
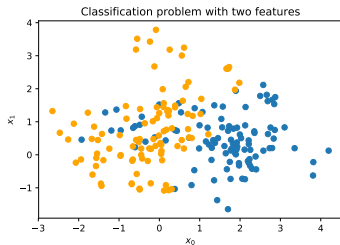
- ▶ This is called “making a prediction” for  $\mathbf{x}_i$ .
- ▶ The entire fitted surface (hyperplane) is fully characterized by the parameter vector  $\hat{\mathbf{w}}$ .
- ▶ After fitting the model, we can “discard” the training dataset.

## Example: Linear model fit by least squares

- ▶ Scatter plot (on next slide) of training data on a pair of inputs  $x_1$  and  $x_2$
- ▶ Output class variable  $g$  has two values **BLUE** and **ORANGE**.
- ▶ Linear regression model fitted with the response variable  $y$  coded as 0 for **BLUE** and 1 for **ORANGE**.
- ▶ Fitted values  $\hat{y}$  converted to a fitted class variable  $\hat{g}$  as

$$\hat{g} = \begin{cases} \text{BLUE} & \text{if } \hat{y} \leq 0.5 \\ \text{ORANGE} & \text{if } \hat{y} > 0.5 \end{cases}$$

# Example: Linear model fit by least squares



## Example: Linear model fit by least squares

- ▶ Two classes separated in the plane ( $\mathbb{R}^2$ ) by the decision boundary  $\{\mathbf{x} : \mathbf{w}^T \mathbf{x} = 0.5\}$
- ▶  $\{\mathbf{x} : \mathbf{w}^T \mathbf{x} < 0.5\}$  set of BLUE points
- ▶  $\{\mathbf{x} : \mathbf{w}^T \mathbf{x} \geq 0.5\}$  set of ORANGE points

## Example: Linear model fit by least squares

- ▶ Wrong classifications on both sides of the boundary
- ▶ Are the errors caused by the model or are they unavoidable?
- ▶ Two possible scenarios
  - ▶ **Scenario 1:** data generated from bivariate Gaussian distribution
  - ▶ **Scenario 2:** data generated from 10 Gaussian distributions; the means of these distributions are also distributed as Gaussian
- ▶ In Scenario 1 the linear boundary is the best we can do since the overlap is inevitable.
- ▶ In Scenario 2 the linear boundary is unlikely to be optimal (in fact the boundary is non-linear and disjoint).

## Discussion point

What is the expression for the Euclidean distance between two vectors (points) **a** and **b** in vector form?

# Nearest-neighbours model

- ▶ In nearest-neighbour methods  $\hat{y}(\mathbf{x})$  is determined based on the inputs (points) in the training set  $\mathcal{T}$  which are "closest" to the input  $\mathbf{x}$ .
- ▶  $k$ -nearest neighbour fit is defined as

$$\hat{y}(\mathbf{x}) = \frac{1}{k} \sum_{\mathbf{x}_i \in N_k(\mathbf{x})} y_i$$

where  $N_k(\mathbf{x})$  is the neighbourhood of  $\mathbf{x}$  consisting of the  $k$  "closest" points to  $\mathbf{x}$ .

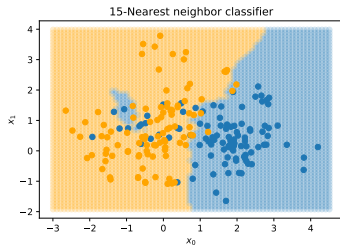
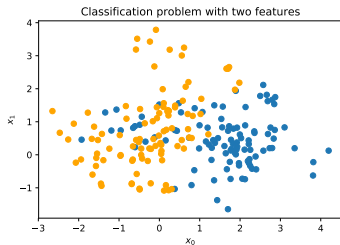
- ▶ "Closeness" requires a definition of **metrics**.
- ▶ For the moment we assume Euclidian distance (each  $\mathbf{x}$  is a point in the hyperspace).
- ▶ An average of the classes of the  $k$  closest points (but only for binary classification problem).

## Back to the BLUE and ORANGE example

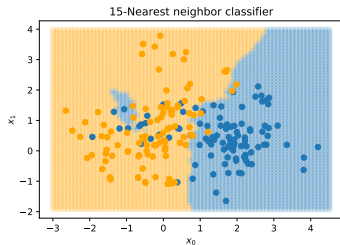
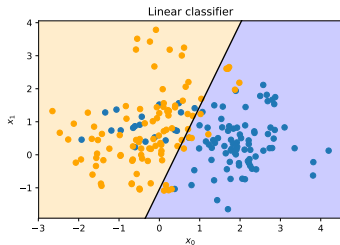
- ▶ We use the same training data as in the linear model example.
- ▶ New borderline between the classes generated with 15-nearest-neighbour model.
- ▶ Since ORANGE is encoded as 1  $\hat{y}$  is the proportion of ORANGE points in the 15-neighbourhood.
- ▶ Class ORANGE assigned to  $\mathbf{x}$  if  $\hat{y}(\mathbf{x}) > 0.5$  (majority is ORANGE).



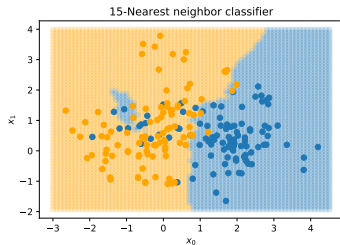
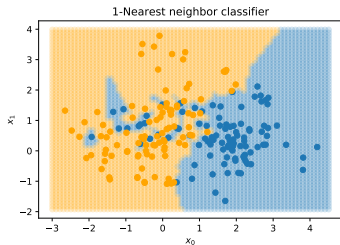
# 15-Nearest neighbour classifier



# Linear classifier vs. 15-Nearest neighbour



# 1-Nearest neighbour vs. 15-Nearest neighbour



# Comparison of the techniques

- ▶ 15-NN seems to work better than the linear classifier since fewer points are missclassified.
- ▶ On the other hand, **none** of the points in the 1-NN case was misclassified!?
- ▶ Actually with the 1-NN method the error on **training data** is always 0.
- ▶ An independent test set needed to obtain a better comparison of the methods.

# Comparison of techniques

- ▶ At first sight it looks like k-NN has only one parameter,  $k$  versus  $p$  parameters (number of weights  $w_i$ ) of the linear model.
- ▶ The **effective** number of parameters of k-NN is  $N/k$  which is in general bigger than  $p$  ( $N$  is the size of the training set).
- ▶ For instance, assume non-overlapping neighbourhoods
  - ▶ There will be  $N/k$  neighbourhoods.
  - ▶ To each neighbourhood there correspond one parameter (the mean of the elements of the neighbourhood).

## Discussion point

Assume that you are building a machine learning model to be used as an *aid* by clinicians for decision making. The inputs to the model are a number of *biomarkers* describing the condition of the patient.

The clinician specifies that they are interested in a model that is *interpretable*, i.e. a model that will not only output a prediction but also give an indication about which biomarkers are important when making the prediction.

You can either use a linear model or a  $k$ -NN classifier. What is the better choice in your opinion?

# Probability theory

Materials:

- ▶ Chapter 1.3 from **deeplearning**

# Probability theory

- ▶ Probability theory is a mathematical framework for dealing with uncertainty, i.e., modeling and analyzing uncertain events and statements
- ▶ In AI probability theory is used in two major ways:
  - ▶ To design AI systems, i.e., derive models and expressions and the corresponding algorithms.
  - ▶ To analyze the behaviour of the AI systems.



# Probability theory

- ▶ A **random variable** is a variable that can take values randomly.
- ▶ We will denote random variables with plain (ordinary text) typeface and their values with standard math typeface for example, if the random variable is denoted as  $x$  its values can be  $x_1$  and  $x_2$ .
- ▶ A vector-valued random variable is denoted with bold typeface, e.g.  $\mathbf{x}$ .
- ▶ On its own a random variable just denotes the set of its possible values; to get its full meaning it needs to be coupled with a distribution.

# Probability theory

- ▶ There are two types of random variables: **discrete** and **continuous**.
- ▶ Consequently there are two ways to describe probability distributions: **probability mass functions** and **probability density functions**.

# Probability mass function

- ▶ The domain of a probability mass function  $P$  is the set of all possible states of the random variable  $x$ .
- ▶  $\forall x \in \mathcal{X} : 0 \leq P(x) \leq 1$ 
  - ▶ An impossible event has probability 0 and no state can be less probable than that.
  - ▶ An event that is guaranteed to happen has probability 1 and no state can have a greater chance of occurring.
- ▶  $\sum_{x \in \mathcal{X}} P(x) = 1$ 
  - ▶ We say that  $x$  is **normalized**.
- ▶ Example: Uniform distribution:  $P(x = x_i) = \frac{1}{k}$ .

# Probability density function

- ▶ The domain of the probability density function  $p$  must be the set of all possible states of  $x$ .
- ▶  $\forall x \in x : p(x) \geq 0$ .



$$\int p(x) dx = 1$$

- ▶ Example: uniform distribution  $u(x; a, b) = \frac{1}{b-a}$ , for  $x \in [a, b]$

# Conditional probability

- ▶ **Conditional probability** is the probability of some event provided that some other event has happened.
- ▶ Given two random variables  $x$  and  $y$ , the conditional probability that  $y$  has value  $y$  provided that we know that  $x$  has value  $x$  is given by

$$P(y = y \mid x = x) = \frac{P(x,y)}{P(x = x)}$$

- ▶ Another way to see this formula is

$$P(x,y) = P(x = x)P(y = y \mid x = x)$$

i.e., the probability of  $x$  and  $y$  occurring together is equal to the probability of occurrence of  $x$  times the probability of  $y$  occurring provided  $x$  has occurred.

# Expectation

- ▶ The **expectation** or **expected** value of a function  $f(x)$  with respect to a probability distribution  $P(x)$  is the average value of  $f$  over all values  $x$  assuming they are drawn from  $P$



$$\mathbb{E}_{x \sim P}[f(x)] = \sum_x P(x)f(x)$$



$$\mathbb{E}_{x \sim P}[f(x)] = \int p(x)f(x)dx$$

- ▶ Linearity of expectations:

$$\mathbb{E}_x[\alpha f(x) + \beta g(x)] = \alpha \mathbb{E}_x[f(x)] + \beta \mathbb{E}_x[g(x)]$$

# Variance and covariance

- ▶ The **variance** gives a measure of variation of the values of a random variable  $x$

$$\text{Var}(f(x)) = \mathbb{E}[(f(x) - E[f(x)])^2]$$

Square root of the variance is called **standard deviation**.

- ▶ The **covariance** is a measure of linear relation as well as scale between

$$\text{Cov}(f(x), g(x)) = \mathbb{E}[(f(x) - E[f(x)])(g(x) - E[g(x)])]$$

# Covariance matrix

- ▶ The **covariance matrix** of a random vector  $\mathbf{x} \in \mathbb{R}^n$  is a  $n \times n$  matrix with elements

$$\text{Cov}(\mathbf{x})_{i,j} = \text{Cov}(x_i, x_j)$$

- ▶ The diagonal elements of the matrix give the variance

$$\text{Cov}(x_i, x_i) = \text{Var}(x_i)$$



# Bernouli Distribution

- ▶ A distribution over a single binary random variable
- ▶ Controlled by a single parameter  $\phi \in [0, 1]$  which corresponds to the probability of the random variable taking the value 1
- ▶ Properties:

$$P(x = 1) = \phi$$

$$P(x = 0) = 1 - \phi$$

$$P(x = x) = \phi^x (1 - \phi)^{1-x}$$

$$\mathbb{E}_x[x] = \phi$$

$$\text{Var}(x) = \phi(1 - \phi)$$

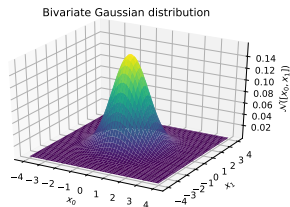
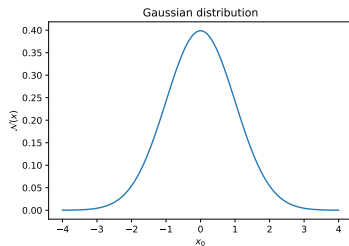
# Gaussian distribution

- ▶ The most commonly used distribution, also called **normal distribution**.
- ▶ Controlled by two parameters  $\mu \in \mathbb{R}$  (the **mean**) and  $\sigma \in (0, \infty)$ , (the **standard deviation**)

$$\mathcal{N}(x; \mu, \sigma^2) = \sqrt{\frac{1}{2\pi\sigma^2}} \exp\left(-\frac{1}{2\sigma^2}(x - \mu)^2\right)$$

$$\mathcal{N}(\mathbf{x}; \boldsymbol{\mu}, \boldsymbol{\Sigma}) = \sqrt{\frac{1}{(2\pi)^n \det(\boldsymbol{\Sigma})}} \exp\left(-\frac{1}{2}(\mathbf{x} - \boldsymbol{\mu})^T \boldsymbol{\Sigma}^{-1}(\mathbf{x} - \boldsymbol{\mu})\right)$$

# Gaussian distribution




# Logistic sigmoid

- ▶ A useful function that we are going to consider

$$\sigma(x) = \frac{1}{1 + \exp(-x)}$$

- ▶ The Logistic (sigmoid) function is commonly used to parametrize Bernoulli distributions.



# Bayes' rule

- ▶ Suppose know  $P(y | x)$ , but we actually need  $P(x | y)$ . If we know  $P(x)$  then we can compute

$$P(x | y) = \frac{P(y | x)P(x)}{P(y)}$$

Although it appears in the formula prior knowledge  $P(y)$  is not needed since usually it can be computed as  $\sum_x P(y | x)P(x)$

- ▶ It can be straightforwardly derived from the conditional probability formula.
- ▶ It could have be named also after Laplace who independently found it, generalized it, and introduced it in practice.

# Acknowledgements

The slides for this lecture were prepared by Mitko Veta and Dragan Bošnjacki.

Some of the slides are based on the accompanying lectures of **deeplearning**.

# References